Graph-based detection of money mules

"What does recent research (2022 to 2026) say about detecting money mule accounts in digital payment systems using graph-based machine learning? Summarise the main methods, their reported accuracy, and their limitations. List the key papers." *Generated by LeapSpace*

Graph-Based Money Mule Detection in Digital Payment Systems (2022-2026): Quick Insights

Key Findings Table

Finding	Evidence base / setting	Reported performance	Main caveat
Graph-based models consistently outperform feature-only baselines because they capture account-to-account structure, shared infrastructure, and fund-flow patterns that tabular models miss.	Broad fraud/AML transaction-network studies; payment and financial graphs [1] [2] [3] [4] [5] [6] [7] [8]	Often better precision/recall/AUC than non-graph baselines, though not always fully quantified in excerpts [9] [10] [11]	Much of the evidence is **fraud/AML-adjacent**, not strictly mule-only [12] [13]
Temporal and dynamic graph learning is the strongest recent trend. Preserving event order and evaluating chronologically appears more realistic than static/random splits.	Temporal GNN, rolling-window, dynamic graph studies [14] [15] [16] [17] [18] [19]	Better fraud metrics with richer event history; temporal validation often reduces optimistic offline scores [18] [20] [21] [22]	Exact gains are often task-specific and not always reported numerically in the supplied extracts [16] [23]
Heterogeneous multi-relational graphs are becoming the dominant industrial design.	Payment graphs with accounts, devices, IPs, merchants, customers, institutions [24] [25] [26] [27] [28]	HT-GNN: **65.8% recall at 0.1% FPR** on 847M transactions / 52M entities [26]; TransFraudGNN: **AUC-ROC 0.9487 / 0.9361 / 0.9218** with **26.1%-34.6% relative AUC error reduction** over strongest baseline [27]	Large-scale industrial results are promising but may not transfer directly across institutions or rails [17] [29]
Mule-specific graph evidence is still thin compared with broader fraud and AML graph research.	Mule-focused items are sparse; strongest direct examples are MuleTrack and MuleTrace [30] [31]	MuleTrack is described as heuristic + Markov-chain temporal modeling for UPI, but standard AUC/F1/recall values were not provided in the retrieved evidence [30]	No complete mule-only benchmark landscape can be reconstructed from the supplied evidence [12] [13]

Finding	Evidence base / setting	Reported performance	Main caveat
Motif, typology, and subgraph methods remain useful, especially where labels are scarce.	Temporal motifs, typology indicators, local neighborhood subgraphs [32] [33] [34]	Typology features increased recall by **9-11 percentage points**; best RF recall **56.3%** in one study [32] [35]	High-recall settings may still create costly false positives [35]
Accuracy is usually the wrong headline metric for mule detection because of severe class imbalance.	Benchmark and review literature [36] [37] [38] [39] [40]	More informative metrics are recall, F1, AUCPR/AUPRC, MCC, G-Mean, and low-FPR operating points [36] [37] [38] [39] [41]	Papers still sometimes report high accuracy without enough minority-class context [36] [42]

Direct Answer

Recent research supports using **graph-based machine learning** for money mule detection in digital payment systems because transaction graphs encode patterns that tabular models miss: intermediary accounts, shared devices/IPs, rapid receive-and-forward behavior, and suspicious neighborhood structure [1] [2] [4] [28]. The strongest methodological trend from 2022-2026 is a shift from **static node classification on simple transaction graphs** toward **temporal, dynamic, and heterogeneous graph models** that better preserve chronology and multi-entity context [14] [15] [16] [17] [26] [27].

However, the literature has an important gap: **few papers directly benchmark labeled money mule accounts**. Most of the quantitative evidence comes from adjacent **fraud detection** or **AML/laundering detection** studies rather than explicit mule-account datasets [12] [13] [30]. So the best current answer is: graph ML appears well suited to mule detection, but the field's strongest numbers mostly come from broader transaction-fraud and AML settings rather than mule-only benchmarks [5] [26] [27] [43].

The main methods and reported results in the supplied evidence are:

- **Static and heterogeneous GNNs**: GraphSAGE, GCN, GAT, GIN, R-GCN, and heterogeneous GNNs for node classification [5] [9] [10] [11] [44] [45].
- **Temporal/dynamic graph models**: TGN/TGAT-style models, LSTM+GCN hybrids, chronological dynamic learning, and Markov-chain hybrids [16] [20] [30] [46].
- **Motif/subgraph/path approaches**: temporal motifs, graph typologies, local neighborhood subgraphs, and contrastive/self-supervised graph learning [32] [33] [34] [47] [48].

The clearest reported numbers in the evidence are:

- **HT-GNN**: **65.8% recall at 0.1% false-positive rate** on a large heterogeneous dynamic payment graph [26].
- **TransFraudGNN**: **AUC-ROC 0.9487, 0.9361, 0.9218** and **26.1%-34.6% relative AUC error reduction** over the strongest baseline [27].

- **Heterogeneous Graph Auto-Encoder for Credit Card Fraud Detection**: **AUC-PR 0.89, F1 0.81** [5].
- **R-GCN drift study**: average **F1 0.868**, more stable than HGT and HAN under drift [43].
- **GraphSAGE Medicare fraud study**: gains over baseline of **+0.01 precision, +0.35 recall, +0.30 F1, +0.18 AUC** [11].
- **GCNShield**: **93.4% accuracy, 96.7% precision**, but these should be interpreted cautiously under imbalance [36] [42].
- **LaundroGraph**: self-supervised link-prediction AUC improved by **12 percentage points** over the best non-graph baseline [47].

Main limitations are consistent across papers: **scarce mule labels, severe class imbalance, concept drift, temporal leakage from random splits, limited external validation, and insufficient interpretability for investigator workflows** [22] [29] [40] [49] [50] [51]. For practical mule detection, the literature implies that systems should optimize for **recall at very low FPR**, use **chronological or rolling validation**, and provide **subgraph-level explanations** rather than relying on headline accuracy alone [18] [26] [39] [52].

Confidence level: High - the overall direction of evidence is consistent across multiple studies and reviews, though **mule-specific quantitative evidence is limited** [12] [13].

Study Scope

- **Time period:** 2022-2026, as defined by the supplied evidence [30] [47] [48].
- **Domains covered:** digital payments, payment rails, AML, transaction fraud, credit-card fraud, Bitcoin/crypto AML, and adjacent financial fraud settings used as methodological evidence [5] [23] [26] [27].
- **Method classes:** static GNNs, heterogeneous GNNs, temporal/dynamic GNNs, motif/subgraph approaches, self-supervised graph learning, and hybrid heuristic + probabilistic systems [11] [20] [30] [34] [45] [47].
- **Tasks covered:** mostly node classification, with some link prediction, anomaly detection, and laundering-chain tracing [31] [47] [53] [54].

Assumptions & Limitations

- The answer necessarily draws on **adjacent fraud/AML graph literature** because direct mule-account benchmarks are sparse [12] [13] [30].
- Some papers report improvements over baselines without enough excerpted detail to quantify every comparison [9] [10] [44].
- High reported accuracy or precision may be misleading under extreme class imbalance unless paired with recall, AUCPR, MCC, or low-FPR analysis [36] [38] [39].
- Several findings come from proprietary or institution-specific data, limiting external validity [26] [29].

Suggested Further Research

- Build **public or reproducible mule-specific graph datasets** with explicit mule labels rather than generic suspicious/fraud labels [13] [55].
- Compare graph models directly against **strong tree ensembles** on the same imbalanced payment data [56] [57].
- Test temporal robustness across **institutions, geographies, and payment rails** using external validation [17] [29].
- Move beyond node classification toward **edge-level and subgraph-level laundering-chain detection** [54] [58].
- Improve **explainability** so analysts receive actionable subgraph narratives rather than opaque risk scores [52] [59].

Full Report

Introduction and Scope

Money mule detection aims to identify accounts that receive and move funds on behalf of others, often acting as intermediaries inside broader laundering or fraud schemes [15] [30] [31]. In digital payment systems, graph-based ML is attractive because payments are inherently relational: suspicious behavior often emerges not from a single transaction feature, but from the **pattern of transfers among accounts and related entities** such as devices, IPs, merchants, or counterparties [1] [2] [3] [28].

The recent literature suggests that graph methods add value precisely where tabular models are weak. They can encode neighborhood structure, repeated interactions, shared infrastructure, and multi-hop fund flows that are central to mule behavior and laundering-chain discovery [1] [2] [4] [28]. Yet the evidence base is uneven: there are many graph-fraud and AML papers, but only a small number directly targeting labeled mule accounts [12] [13] [30].

Confidence level: High - this scope statement is directly supported by multiple cross-cutting findings [1] [2] [3] [13].

Foundational Context: Why Graphs Matter for Mule Detection

Three concepts are necessary to interpret the recent literature.

1) Transaction graphs as the core data model

Most studies construct a graph where **accounts or users are nodes** and **transactions are directed edges**, often with amount, timestamp, and direction as edge attributes [24] [60] [61]. Some papers expand this into **heterogeneous graphs** containing users, transactions, merchants, devices, cards, IP addresses, or institutions as separate node or relation types [25] [26] [27].

2) Mule detection is usually framed as node classification

Most recent graph-fraud papers cast the task as predicting whether an account, address, or transaction node is fraudulent or suspicious [53] [62] [63] [64] [65]. This is practical, but it only partially matches the laundering problem, because mule activity also has **edge-level** and **subgraph-level** structure, such as bursty in-out chains or ring-like flows [31] [54] [58].

3) Imbalance and drift are first-order constraints

Mule accounts are rare, labels are sparse, and fraud tactics evolve. The literature therefore argues that **accuracy alone is insufficient**, and that evaluation should prioritize recall, precision-recall behavior, MCC, G-Mean, and low-FPR operating points under chronological testing [18] [36] [37] [38] [39].

Taken together, these findings suggest that a useful mule-detection system is not just "a GNN on transactions"; it is typically a **time-aware, imbalance-aware, multi-relational graph pipeline** [17] [26] [40].

Confidence level: High - the methodological foundations recur across several strands of the evidence [24] [39] [64].

How Transaction Graphs Are Built

Nodes, edges, and attributes

Recent studies are mostly transaction-centric. The basic pattern is accounts/users as nodes with directed payment edges and attributes such as timestamp, amount, and direction [24] [60] [61]. More advanced systems add devices, IP addresses, merchants, cards, institutions, or transaction records as additional entity types [25] [26] [27] [66].

This design matters because mule behavior often spans entities beyond the account itself. Shared devices or IPs can connect apparently unrelated accounts, and merchant or institution links can expose organized structures that a flat account-to-account graph would miss [26] [27] [28].

Time-resolved and dynamic design

Several studies emphasize preserving chronology. Dynamic or temporal graphs capture changing behavior and reduce leakage from future information [14] [15] [18] [67]. Some use explicit time windows such as recent months to control graph size, while others allow multiple directed edges between the same accounts to preserve event history [24] [60] [67].

Design choices reported to improve detection

The strongest recurring design choices are:

- **heterogeneous/multi-relational graphs** [26] [27] [68]

- **dynamic or temporal modeling** [16] [20] [67]
- **attention mechanisms** to emphasize informative neighbors or relations [69] [70]
- **motif and typology features** to encode suspicious local structures [32] [34]
- **local neighborhood subgraphs** rather than full-network processing in some settings [33]

Several sources explicitly state that graph-based representations outperform tabular baselines because they preserve complex fund-flow structure across interacting entities [69] [71] [72].

Confidence level: High - graph construction patterns are consistent across industrial, review, and methodological sources [24] [26] [28].

Main Method Families and Reported Performance

1) Static and heterogeneous GNNs

The literature repeatedly cites GCN, GraphSAGE, GAT, GIN, R-GCN, and heterogeneous GNN variants as the main supervised model families for transaction-graph fraud detection [5] [9] [10] [11] [44] [45].

Reported results include:

- **GraphSAGE Medicare fraud detection (2022)**: better than baseline by **+0.01 precision, +0.35 recall, +0.30 F1, +0.18 AUC** [11].
- **Heterogeneous Graph Auto-Encoder for Credit Card Fraud Detection (2025)**: **AUC-PR 0.89, F1 0.81** [5].
- **GCNShield (2026)**: **93.4% accuracy, 96.7% precision** [42].
- **R-GCN drift study (2025)**: average **F1 0.868**, with R-GCN more robust than HGT and HAN; HGT reportedly dropped from **0.747 to 0.455** under drift [43].

The caution is that some studies report strong headline scores without enough context on class balance or deployment realism [36] [42].

2) Heterogeneous industrial-scale payment models

This is the clearest high-performance strand for payment systems specifically.

- **HT-GNN** models cards, devices, merchants, customers, and IPs in a dynamic heterogeneous graph over **847M transactions and 52M entities**, reporting **65.8% recall at 0.1% FPR**; ablations attribute gains partly to heterogeneity, device/IP relations, and temporal attention [26].
- **TransFraudGNN** uses heterogeneous interaction graphs with users, merchants, devices, and behavior entities, reporting **AUC-ROC 0.9487, 0.9361, 0.9218** and **26.1%-34.6% relative AUC error reduction** over the strongest baseline [27].

- **xFraud** is presented as an explainable heterogeneous GNN for industrial-scale transaction fraud, paired with a hybrid explainer for analyst review [52].

These results support the broader meta-pattern that industrial payment systems are moving toward **multi-entity, relation-specific graph learning** [26] [27] [52].

3) Self-supervised and AML-oriented graph systems

- **LaundroGraph (2022)** models AML as a customer-transaction bipartite graph and uses self-supervised GNN link prediction, achieving a **12 percentage-point AUC improvement** over the best non-graph baseline [47].
- **Inspection-L (2023)** combines DGI-style self-supervision with a GIN encoder and supervised RF for illicit transaction detection on Elliptic, reporting state-of-the-art performance on key classification metrics in the provided summary, though exact numbers are not included here [48].

These systems are important because they address one of the field's hardest constraints: **label scarcity** [47] [49].

4) Motif, typology, and local subgraph methods

Where explicit mule labels are limited, some studies use engineered graph structure.

- Graph-based typology indicators increased recall by **9-11 percentage points**, with best RF recall **56.3%** [32] [35].
- Temporal motifs outperformed simpler graph features and embedding methods on fraud tasks in Mercari and JPMC payment networks [34].
- A **3-hop neighborhood subgraph** approach with GCN+LSTM outperformed other GCN-based methods for illicit transaction detection while using only about **80 nodes on average** rather than the whole network [33].

These approaches are often easier to align with specific laundering typologies, but high recall can come with high false-positive cost [35].

Confidence level: Medium-High - the model families and headline metrics are well supported, but some excerpts omit exact baseline numbers [9] [10] [44].

Temporal and Dynamic Graph Models

The most consistent cross-paper trend is the move from static graphs to temporal and dynamic learning [14] [15] [16] [17].

Why temporal modeling helps

Temporal graph work argues that snapshot methods are sensitive to window choice, while event-based models such as TGN-style architectures preserve the order and timing of interactions through time encoding and recurrent state updates [16]. This matters because fraud and mule behavior is often bursty and non-stationary [17] [73].

Reported temporal evidence

- TGN-based work on online payments found that combining richer event histories improved fraud metrics compared with single-event graphs [20].
- Another temporal-graph payment study also reported better AUC/recall when multiple events were combined [21].
- LSTM+GCN temporal hybrids were reported as successful on Bitcoin or card-related fraud tasks [46] [74].
- TAGA, based on TGAT-style temporal encoding, is presented as promising for real-time AML monitoring [23].
- A deployed-style LSTM/GNN/federated-XAI architecture reported **latency below 110 ms per transaction stream** and **retraining under 1.9 s per adaptive window**, supporting streaming use [75].

Evaluation under drift

Recent papers increasingly use **chronological and rolling validation** to mimic deployment and avoid leakage [18] [19]. DISCO uses ten rolling temporal splits, and SCALE uses chronological 6/2/2 windows specifically to prevent temporal leakage and test unseen fraud patterns [18] [19]. Other work explicitly warns that offline retrospective metrics do not capture latency, streaming operations, or analyst workflow integration [22].

Taken together, these findings suggest that temporal modeling is not just a performance tweak; it is a **validity requirement** for operational mule detection [17] [18].

Confidence level: High - this is one of the most consistent findings in the entire evidence base [14] [16] [18] [19].

Heterogeneous and Multi-Relational Graph Models

Heterogeneous graph design is increasingly central in payment ecosystems [25] [26] [27].

Why heterogeneity helps

A mule ring may be weakly connected in pure account-transfer space but obvious once shared devices, IPs, merchants, or institutions are included [26] [28]. Relation-specific message passing can preserve the meaning of each link type rather than collapsing them into one undifferentiated edge set [45] [76].

Explainability and analyst support

Recent industrial work increasingly treats interpretability as operationally necessary rather than optional [52] [59] [77]. xFraud combines a heterogeneous detector with a hybrid explainer using GNNExplainer-style and centrality-based signals to support review [52]. More general XGNN-style methods extract explanatory subgraphs or counterfactuals to show what drove a prediction [59].

Foundational methods reused in AML/fraud

Key foundations repeatedly referenced include:

- **R-GCN** for multi-relational reasoning [45]
- **HAN** for hierarchical attention over heterogeneous relations [76]
- **GraphSAGE** for inductive scoring of unseen nodes [78]
- **CARE-GNN** for camouflage-aware neighbor filtering [79]
- **PC-GNN** for class-imbalance-aware sampling [80]

These foundations matter because production payment systems continuously score new accounts and must handle multiple relation types under extreme imbalance [78] [80].

Confidence level: High - heterogeneous graph modeling is directly supported by both industrial systems and foundational method links [26] [27] [45].

Evaluation Practice, Accuracy Claims, and Recurring Limitations

Metrics: what should matter

The literature is clear that **accuracy is insufficient** for rare mule/fraud classes [36] [38] [39]. More informative metrics are:

- **Recall** for finding rare mules [36] [39]
- **Precision / low-FPR operating points** for investigator burden [26] [39]
- **F1** for a thresholded balance [37] [38]
- **AUCPR/AUPRC** because they are less inflated by true negatives [36] [37]
- **MCC and G-Mean** as imbalance-aware summaries [38] [39]

Class imbalance and label scarcity

Most retrieved studies frame the task as node classification but note that severe minority scarcity is the main bottleneck [40] [64]. Common responses include resampling, SMOTE, undersampling, class-balanced or focal losses, and sampling strategies tailored to graph neighborhoods [80] [81] [82] [83] [84] [85].

Self-supervised graph learning is especially relevant because it reduces dependence on scarce labels [47] [48].

External validity and deployment realism

The main reported limitations are:

- **sparse or proxy labels** [40] [49]
- **concept drift and non-stationarity** [17] [50]
- **temporal leakage from random splits** [18] [86]
- **single-dataset or proprietary evaluation**, limiting transfer [29] [51]
- **black-box behavior and limited interpretability** [40] [59]
- **few public mule-specific datasets** [13] [55]

This creates a recurring contradiction: some papers report very high accuracy or precision, but operationally useful systems must prove performance under drift, low FPR, delayed labels, and human investigation constraints [22] [39] [42].

Confidence level: High - limitations are one of the most repeatedly cited themes across reviews and benchmark discussions [29] [40] [50].

Key Papers to Highlight

Most directly relevant mule / AML / payment-graph papers

- **MuleTrack (2026)** - direct mule detection in India's UPI using heuristics plus Markov-chain temporal modeling [30].
- **MuleTrace** - uses money mules as leads for unsupervised laundering-chain detection [31].
- **NPCI billion-scale anomaly detection (2026)** - large UPI account anomaly setting with 400M accounts over 3 months, though labels are suspicious/fraudulent rather than explicit mule labels [55].
- **HT-GNN** - dynamic heterogeneous graph for payment fraud at very large scale [26].
- **TransFraudGNN** - heterogeneous interaction graph with strong AUC results [27].
- **xFraud** - explainable heterogeneous GNN for industrial transaction fraud [52].
- **LaundroGraph (2022)** - self-supervised AML graph learning [47].
- **Inspection-L (2023)** - self-supervised + GIN + RF for illicit transaction detection [48].

Important benchmark / baseline papers

- **CARE-GNN** [79]
- **PC-GNN** [80]

- **GraphConsis** [87]
- **R-GCN** [45]
- **HAN** [76]
- **GraphSAGE** [78]

Useful recent quantitative comparators

- **Heterogeneous Graph Auto-Encoder for Credit Card Fraud Detection (2025)** [5]
- **Semantic and Structural Drift in Financial Knowledge Graphs (2025)** [43]
- **GCNShield (2026)** [42]

Conclusion

Recent research indicates that graph-based ML is a strong fit for detecting money mule accounts in digital payment systems, but the empirical base is still dominated by **broader fraud and AML studies rather than mule-only benchmarks** [12] [13]. The field is moving toward **heterogeneous, temporal, and dynamic graph models** because they better represent real payment ecosystems and evolving suspicious behavior [14] [17] [26] [27]. The best reported results in the supplied evidence come from such systems, including **HT-GNN's 65.8% recall at 0.1% FPR**, **TransFraudGNN's AUC-ROC up to 0.9487**, and **graph auto-encoder performance of AUC-PR 0.89 / F1 0.81** [5] [26] [27].

The practical message is narrower than the headline numbers: for real mule detection, **chronological validation, low-FPR operation, label-efficient training, and explainable subgraph outputs** are at least as important as raw offline accuracy [18] [39] [52]. The main unresolved research problem is not whether graphs help—they do—but whether the field can produce **mule-specific, temporally realistic, externally validated benchmarks** that prove these methods work robustly across live payment systems [13] [29] [55].

Overall confidence level: High - the broad methodological direction and main limitations are strongly supported across the supplied evidence; **confidence is lower only for mule-specific quantitative comparisons because the direct benchmark literature is sparse** [12] [13] [30].

References

1. Balaji, E.N., Ganesh, P., Sindhuja, M. (2025) Improved Detection of Anti-Money Laundering in Bank Transactions Using Graph Neural Networks. *2025 International Conference on Sustainability, Innovation and Technology, ICSIT 2025*. https://doi.org/10.1109/ICSIT65336.2025.11294874
2. Sathyamoorthy, P., Lenin Raj, S.J., Nagarajan, N. (2026) Multi-Perspective Financial Fraud Detection Using Integrated Learning Models and Cloud Computing Platforms. *Proceedings of 8th International Conference on Intelligent Sustainable Systems, ICISS 2026*. https://doi.org/10.1109/ICISS67859.2026.11453893

3. Bukhori, H.A., Munir, R. (2023) Inductive Link Prediction Banking Fraud Detection System Using Homogeneous Graph-Based Machine Learning Model. *2023 IEEE 13th Annual Computing and Communication Workshop and Conference, CCWC 2023*. https://doi.org/10.1109/CCWC57344.2023.10099180
4. Sahith, Y., Reddy, K.S.S.T., Ganesh, V.J., (...), Rashmi, M.R. (2026) Detecting Fraudulent Transactions: Leveraging Graph Neural Networks for Enhanced Probability Analysis. *2026 International Conference on Smart Futuristic Technology, ICSFT 2026*. https://doi.org/10.1109/ICSFT66733.2026.11507561
5. Majumder, S., Singh, M.T., Prasad, R.K., (...), Singh, N. (2025) Heterogeneous Graph Auto-Encoder for Credit Card Fraud Detection. *International Journal of Computers and their Applications*. https://www.scopus.com/pages/publications/105009616291
6. Liu, W., Hallaji, E., Razavi-Far, R. (2026) Graph Neural Networks for Class Imbalance Learning and Detecting Fraud Transactions. *2026 IEEE Conference on Artificial Intelligence, CAI 2026*. https://doi.org/10.1109/CAI68641.2026.11536192
7. Sultana, I., Maheen, S.M., Kshetri, N., Fardous Zim, M.N. (2025) detectGNN: Harnessing Graph Neural Networks for Enhanced Fraud Detection in Credit Card Transactions. *ISDFS 2025 - 13th International Symposium on Digital Forensics and Security*. https://doi.org/10.1109/ISDFS65363.2025.11011957
8. Thilagavathi, M., Saranyadevi, R., Vijayakumar, N., (...), K, K. (2024) AI-Driven Fraud Detection in Financial Transactions with Graph Neural Networks and Anomaly Detection. *Proceedings of 2024 International Conference on Science, Technology, Engineering and Management, ICSTEM 2024*. https://doi.org/10.1109/ICSTEM61137.2024.10560838
9. Rakesh, M., Shetty D, D. (2024) Enhanced Medicare Fraud Detection Using Graph Convolutional Networks. *2024 4th International Conference on Artificial Intelligence and Signal Processing, AISP 2024*. https://doi.org/10.1109/AISP61711.2024.10870744
10. Lu, S. (2025) Graph Neural Network Model in Financial Fraud Detection. *Proceedings - 2025 2nd International Conference on Intelligent Computing and Robotics, ICICR 2025*. https://doi.org/10.1109/ICICR65456.2025.00176
11. Yoo, Y., Shin, D., Han, D., (...), Shin, J. (2022) Medicare fraud detection using graph neural networks. *International Conference on Electrical, Computer, and Energy Technologies, ICECET 2022*. https://doi.org/10.1109/ICECET55527.2022.9872963
12. Verlaan, S., HulpuÈ™, I., van Leeuwen, E.J. (2025) LAS-GNN: A Graph Neural Network for Temporal Money Laundering Motif Detection. *ICAIF 2025 - 6th ACM International Conference on AI in Finance*. https://doi.org/10.1145/3768292.3770410
13. Rawat, S.S., Mishra, A. (2023) The Best ML Classifier(s): An empirical study on the learning of imbalanced and resampled credit card data. *Proceedings of 2023 2nd International Conference on Informatics, ICI 2023*. https://doi.org/10.1109/ICI60088.2023.10421691

14. Sizan, M.M.H., Chouksey, A., Dola, A., (...), Oishi, R.A. (2025) MACHINE LEARNING-BASED UNSUPERVISED ENSEMBLE APPROACH FOR DETECTING NEW MONEY LAUNDERING TYPOLOGIES IN TRANSACTION GRAPHS. *International Journal of Applied Mathematics*. https://doi.org/10.12732/ijam.v38i2s.88

15. Tariq, H., Hassani, M. (2025) Topology-Agnostic Detection of Temporal Money Laundering Flows in Billion-Scale Transactions. *Communications in Computer and Information Science*. https://doi.org/10.1007/978-3-031-74643-7_29

16. Zhen, W., Dai, Q., Chen, L., Liu, D. (2026) TF-GNN: A temporal feedback graph neural network with discriminative regularization for heterogeneous fraud detection. *Neurocomputing*. https://doi.org/10.1016/j.neucom.2026.133694

17. https://www.sciencedirect.com/science/article/pii/S2364416826003513

18. Zhu, R., Zhang, X., Wang, T., (...), Zhang, X. (2026) DISCO: Decoupling representation learning and risk control for reliable credit card fraud detection. *Decision Support Systems*. https://doi.org/10.1016/j.dss.2026.114717

19. Du, J., Hu, C., Zhao, H., (...), Zhao, J. (2026) SCALE: Adaptive heterogeneous graph prototype learning for open-set unknown financial fraud detection. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2026.132749

20. Saldana-Ulloa, D., De-Ita-Luna, G., Marcial-Romero, J.R. (2024) A Temporal Graph Network Algorithm for Detecting Fraudulent Transactions on Online Payment Platforms. *Algorithms*. https://doi.org/10.3390/a17120552

21. Saldana-Ulloa, D., De-Ita-Luna, G. (2025) Fraudulent Event Detection via Temporal Graph Networks. *Proceedings - 2025 IEEE Conference on Artificial Intelligence, CAI 2025*. https://doi.org/10.1109/CAI64502.2025.00168

22. Anh Luong, N.D., Xie, S. (2026) Explainable ensemble machine learning for financial transaction fraud detection: Insights from XGBoost and deep neural networks. *Journal of Finance and Data Science*. https://doi.org/10.1016/j.jfds.2026.100195

23. Zhang, Q., Zhu, Y., Zhang, R., (...), Lan, T. (2025) Research on Anti-Money Laundering Technology Based on Graph Attention Mechanism. *Proceedings of SPIE - The International Society for Optical Engineering*. https://doi.org/10.1117/12.3056070

24. Wang, L., Xu, M., Cheng, H. (2023) Phishing scams detection via temporal graph attention network in Ethereum. *Information Processing and Management*. https://doi.org/10.1016/j.ipm.2023.103412

25. Motie, S., Raahemi, B. (2024) Financial fraud detection using graph neural networks: A systematic review. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2023.122156

26. Chilakalapalli, U.K., Sachdev, P., Mohan, B., Khandelwal, A. (2026) Graph Neural Networks for Real-Time Transaction Fraud Detection in Payment Networks. *2026 International Conference on Smart Applications, Communications and Networking, SmartNets 2026*. https://doi.org/10.1109/SmartNets69662.2026.11604901

27. Zheng, M. (2026) Causal–guided dual–timescale graph neural networks for explainable fraud detection in mobile payment systems. *Journal of King Saud University - Computer and Information Sciences*. https://doi.org/10.1007/s44443-026-00950-w

28. Ghosh, S., Anand, R., Bhowmik, T., Chandrashekhar, S. (2023) GoSage: Heterogeneous Graph Neural Network Using Hierarchical Attention for Collusion Fraud Detection. *ICAIF 2023 - 4th ACM International Conference on AI in Finance*. https://doi.org/10.1145/3604237.3626856

29. https://www.sciencedirect.com/science/article/pii/S2661890725003308

30. Jambhrunkar, G., Sharma, H., Singla, S., Kailasam, T. (2026) MuleTrack: A Lightweight Temporal Learning Framework for Money Mule Detection in Digital Payments. *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-032-02725-2_30

31. Huang, Z. (2025) Enhancing Anti-Money Laundering by Money Mules Detection on Transaction Graphs. *Proceedings of 2025 International Conference on Generative Artificial Intelligence for Business, GAIB 2025*. https://doi.org/10.1145/3766918.3766933

32. Rakhmetulaeva, S., Kulbayeva, A., Bolshibayeva, A., Serbin, V. (2025) IDENTIFYING THE GRAPH-BASED TYPOLOGY FEATURES FOR MACHINE LEARNING MODELS IN FINANCIAL FRAUD DETECTION. *Eastern-European Journal of Enterprise Technologies*. https://doi.org/10.15587/1729-4061.2025.327410

33. Jin, S., Yang, Q., Chen, S., Zhang, H. (2024) Neighborhood Subgraph-based Illicit Transaction Detection in Cryptocurrency Networks. *Proceedings - 2024 International Conference on Identification, Information and Knowledge in the Internet of Things, IIKI 2024*. https://doi.org/10.1109/IIKI65561.2024.00028

34. Liu, P., Altun, B., Acharyya, R., (...), Sariyüce, A.E. (2026) Temporal Motifs for Financial Networks: A Study on Mercari, JPMC, and Venmo Platforms. *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-032-13513-1_18

35. Yiming Lei, Qiannan Shen, Junhao Song (2026) A multi-scale graph learning framework with temporal consistency constraints for financial fraud detection in transaction networks under non-stationary conditions. *International Journal of Data Science and Analytics*. https://doi.org/10.1007/s41060-026-01256-w

36. Sowan, B. (2026) GAKEA-Shield: Knowledge-augmented graph attention over reconstructed attribute-proximity graphs for bitcoin ransomware detection. *Results in Engineering*. https://doi.org/10.1016/j.rineng.2026.112554

37. Zhang, J., Fang, Z., Tang, Y., Ni, L. (2026) TS-GNN: Modeling temporal-semantic discrepancy for graph-based credit card fraud detection. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2026.133233

38. Yu, C., He, X., Wang, J., Zhang, B. (2026) GSLA: Graph fraud detection with structure enhancement and label augmentation under limited supervision. *Neurocomputing*. https://doi.org/10.1016/j.neucom.2025.132291

39. Zhang, J., Wang, T., Ng, W.W.Y., Pedrycz, W. (2022) Ensembling perturbation-based oversamplers for imbalanced datasets. *Neurocomputing*. https://doi.org/10.1016/j.neucom.2022.01.049

40. Manurung, H., Arinengsih, D.R. (2026) Robust machine learning for fraud detection under imbalance, rarity, and adversarial behavior: Evidence across financial reporting, payments, and graph ecosystems. *Social Sciences and Humanities Open*. https://doi.org/10.1016/j.ssaho.2026.103697

41. Zhu, R., Zhang, X., Wang, T., (...), Zhang, X. (2026) DISCO: Decoupling representation learning and risk control for reliable credit card fraud detection. *Decision Support Systems*. https://doi.org/10.1016/j.dss.2026.114717

42. Patel, A., Kumar, A., Singh, S.P. (2026) GCNShield: Advanced Fraud Detection in Payment Systems Using Graph Convolutional Networks. *2026 International Conference on Intelligent Systems in Engineering, Secured Systems and Cybersecurity, ICISESSC 2026*. https://doi.org/10.1109/ICISESSC68634.2026.11542599

43. Menezes, R.S., Holanda Filho, R.H. (2025) Semantic and Structural Drift in Financial Knowledge Graphs: A Robustness Analysis of GNN-Based Fraud Detectors. *Proceedings - IEEE International Conference on Knowledge Graph, ICKG 2025*. https://doi.org/10.1109/ICKG66886.2025.00044

44. Sha, Q., Tang, T., Du, X., (...), Sheng, Y. (2025) Detecting Credit Card Fraud via Heterogeneous Graph Neural Networks with Graph Attention. *2025 IEEE 6th International Seminar on Artificial Intelligence, Networking and Information Technology, AINIT 2025*. https://doi.org/10.1109/AINIT65432.2025.11035158

45. Schlichtkrull, M., Kipf, T.N., Bloem, P., (...), Welling, M. (2018) Modeling Relational Data with Graph Convolutional Networks. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. https://doi.org/10.1007/978-3-319-93417-4_38

46. Sindhu, C., Shivany, R.K., Barakkath Nisha, B.U., Pradeep, G. (2025) Anti-Money Laundering in Bitcoin using Deep Learning. *16th International Conference on Advances in Computing, Control, and Telecommunication Technologies, ACT 2025*. https://www.scopus.com/pages/publications/105034832493

47. Cardoso, M., Saleiro, P., Bizarro, P. (2022) LaundroGraph: Self-Supervised Graph Representation Learning for Anti-Money Laundering. *Proceedings of the 3rd ACM International Conference on AI in Finance, ICAIF 2022*. https://doi.org/10.1145/3533271.3561727

48. Lo, W.W., Kulatilleke, G.K., Sarhan, M., (...), Portmann, M. (2023) Inspection-L: self-supervised GNN node embeddings for money laundering detection in bitcoin. *Applied Intelligence*. https://doi.org/10.1007/s10489-023-04504-9

49. Badjie, A., Xia, Q., Gao, J., (...), Doe, G. (2026) MGCNLLaMA: Multi-graph convolutional networks with lightweight LLaMA-style transformer encoder integration for dynamic fraud detection in cryptocurrency networks. *Journal of Information Security and Applications*. https://doi.org/10.1016/j.jisa.2026.104633

50. Gamal, N., Younis, E.M.G., Mohamed, W.M. (2026) An explainable ensemble learning approach for financial fraud detection: Integrating SHAP and LIME on the PaySim dataset. *Information Processing and Management*. https://doi.org/10.1016/j.ipm.2026.104967
51. Sowan, B. (2026) GAKEA-Shield: Knowledge-augmented graph attention over reconstructed attribute-proximity graphs for bitcoin ransomware detection. *Results in Engineering*. https://doi.org/10.1016/j.rineng.2026.112554
52. https://www.sciencedirect.com/science/article/pii/S2731080925001693
53. Hou, X., Liu, Y., Li, Q., (...), Gao, J. (2026) SRAF: A structure-aware relation-guided model for fraud detection on heterogeneous graphs. *Applied Soft Computing*. https://doi.org/10.1016/j.asoc.2026.115108
54. Capuano, N., Carletti, V., Foggia, P., (...), Vento, M. (2026) Graph neural networks for IoT security: A comparative study. *Internet of Things (The Netherlands)*. https://doi.org/10.1016/j.iot.2025.101863
55. Dasari, B., Dhiraj, T.S., Jambhrunkar, G., (...), Simmhan, Y. (2026) Billion-Scale Fintech Analytics: Scalable Data Management and Anomaly Detection at NPCI. *Proceedings - 2026 IEEE 42nd International Conference on Data Engineering, ICDE 2026*. https://doi.org/10.1109/ICDE65706.2026.00300
56. Waleed, M., Iqbal, N., Alsafi, T. (2026) Bayesian-Optimized Gradient Boosting for Credit Card Fraud Detection: A Cost-Sensitive Evaluation Under Hybrid Resampling. *Lecture Notes in Networks and Systems*. https://doi.org/10.1007/978-3-032-31755-1_16
57. Azbergen, T., Tleubayeva, A., Aituov, A. (2026) Calibration-Aware Gradient Boosting with Temporal Behavioral Features for Anti-Money Laundering Detection. *SIST 2026 - 2026 IEEE 6th International Conference on Smart Information Systems and Technologies, Conference Proceedings*. https://doi.org/10.1109/SIST61674.2026.11596393
58. Ye, M., Han, D., Tang, M., (...), Castiglione, A. (2026) MVHGF: A multi-view hypergraph fusion framework for detecting phishing scams on Ethereum. *Applied Soft Computing*. https://doi.org/10.1016/j.asoc.2026.115071
59. Yagin, Fatma Hilal, Pinar, Abdulvahap (2025) Explainable artificial intelligence in medical research: A synopsis for clinical practitioners—Comprehensive XAI methodologies. *Explainable AI in Healthcare Imaging for Medical Diagnoses*. https://doi.org/10.1016/B978-0-443-23979-3.00012-9
60. Wu, B., Chao, K.-M., Li, Y. (2024) Heterogeneous graph neural networks for fraud detection and explanation in supply chain finance. *Information Systems*. https://doi.org/10.1016/j.is.2023.102335
61. Lin, W., Shi, Y., Liu, Z. (2026) Dynamic multi-level graph enhanced contrastive learning for financial fraud detection. *Pacific Basin Finance Journal*. https://doi.org/10.1016/j.pacfin.2026.103172
62. Zhang, J., Fang, Z., Tang, Y., Ni, L. (2026) TS-GNN: Modeling temporal-semantic discrepancy for graph-based credit card fraud detection. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2026.133233

63. Wang, L., Xu, M., Cheng, H. (2023) Phishing scams detection via temporal graph attention network in Ethereum. *Information Processing and Management*. https://doi.org/10.1016/j.ipm.2023.103412

64. Tong, G., Shen, J. (2023) Financial transaction fraud detector based on imbalance learning and graph neural network. *Applied Soft Computing*. https://doi.org/10.1016/j.asoc.2023.110984

65. Hu, Xinxin, Chen, Haotian, Chen, Hongchang, (...), Liu, Shuxin (2022) Mining Mobile Network Fraudsters with Augmented Graph Neural Networks. *Entropy*. https://doi.org/10.3390/e25010150

66. Manne, V.T. (2026) Adaptive Graph-Based Risk Scoring for Real-Time Instant Payment Systems. *2026 2nd International Conference on Intelligent Systems for Communication, IoT and Security, ICISCoIS 2026*. https://doi.org/10.1109/ICISCoIS62701.2026.11447558

67. Badjie, A., Ntuala, G.M., Xia, Q., (...), Xia, H. (2025) MGGPT: A Multi-Graph GPT-enhanced framework for dynamic fraud detection in cryptocurrency networks. *Computer Networks*. https://doi.org/10.1016/j.comnet.2025.111508

68. https://www.sciencedirect.com/science/article/pii/S2198605325003335

69. Sowan, B. (2026) GAKEA-Shield: Knowledge-augmented graph attention over reconstructed attribute-proximity graphs for bitcoin ransomware detection. *Results in Engineering*. https://doi.org/10.1016/j.rineng.2026.112554

70. Du, J., Hu, C., Zhao, H., (...), Zhao, J. (2026) SCALE: Adaptive heterogeneous graph prototype learning for open-set unknown financial fraud detection. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2026.132749

71. https://www.sciencedirect.com/science/article/pii/S2364416825001659

72. Mukhamadiyev, Abdinabi, Nazarov, Fayzullo, Yarmatov, Sherzod, Cho, Jinsoo (2025) Intelligent Algorithms for the Detection of Suspicious Transactions in Payment Data Management Systems Based on LSTM Neural Networks. *Sensors (Basel, Switzerland)*. https://doi.org/10.3390/s25216683

73. Du, J., Hu, C., Zhao, H., (...), Zhao, J. (2026) SCALE: Adaptive heterogeneous graph prototype learning for open-set unknown financial fraud detection. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2026.132749

74. Dai, S. (2025) Research on Credit Card Risk Prediction Using an LSTM-TCN Hybrid Model. *2025 IEEE 3rd International Conference on Electrical, Automation and Computer Engineering, ICEACE 2025*. https://doi.org/10.1109/ICEACE67491.2025.11439090

75. Al-Dulaimi, H.A., Furaijl, H.B., Baddour, L.S., (...), Ibrahem, D.K. (2025) AI-Driven Behavioral Anomaly and Fraud Detection Models for Real-Time High-Frequency Financial Transactions in FinTech Systems. *ICCR 2025 - 3rd International Conference on Cyber Resilience*. https://doi.org/10.1109/ICCR67387.2025.11292440

76. Wang, X., Ji, H., Cui, P., (...), Ye, Y. (2019) Heterogeneous graph attention network. *The Web Conference 2019 - Proceedings of the World Wide Web Conference, WWW 2019*. https://doi.org/10.1145/3308558.3313562

77. Motie, S., Raahemi, B. (2024) Financial fraud detection using graph neural networks: A systematic review. *Expert Systems with Applications*. https://doi.org/10.1016/j.eswa.2023.122156

78. Hamilton, W.L., Ying, R., Leskovec, J. (2017) Inductive representation learning on large graphs. *Advances in Neural Information Processing Systems*. https://www.scopus.com/pages/publications/85046897776

79. Dou, Y., Liu, Z., Sun, L., (...), Yu, P.S. (2020) Enhancing Graph Neural Network-based Fraud Detectors against Camouflaged Fraudsters. *International Conference on Information and Knowledge Management, Proceedings*. https://doi.org/10.1145/3340531.3411903

80. Liu, Y., Ao, X., Qin, Z., (...), He, Q. (2021) Pick and choose: A GNN-based imbalanced learning approach for fraud detection. *The Web Conference 2021 - Proceedings of the World Wide Web Conference, WWW 2021*. https://doi.org/10.1145/3442381.3449989

81. Baruah, A., Sambhrama, S., Kaliraj, S. (2026) Performance Analysis of a Scalable Machine Learning Model for Financial Fraud Detection. *4th International Conference on Sustainable Computing and Smart Systems, ICSCSS 2026 - Proceedings*. https://doi.org/10.1109/ICSCSS69635.2026.11645977

82. Farrukh, H., Zafar, S., Rehman, Z.U., (...), Alshammry, N. (2025) Blockchain-Based Fraud Detection: A Comparative Systematic Literature Review of Federated Learning and Machine Learning Approaches. *Electronics (Switzerland)*. https://doi.org/10.3390/electronics14244952

83. Aziziyah, R.N., Prasetya, P.P., Bazry, Z.F., Puspita, R. (2025) A Comparative Study of Machine Learning Models for Online Payment Fraud Detection. *2025 International Conference on Computing and Applied Informatics, ICCAI 2025*. https://doi.org/10.1109/ICCAI65301.2025.11279321

84. Oluwadele, D., Malusi, S. (2025) Performance Analysis of ML Algorithms for Fraud Detection in Digital Financial Transactions. *2025 Annual IEEE Conference on Information Communication Technology and Society, ICTAS 2025 - Proceedings*. https://doi.org/10.1109/ICTAS64866.2025.11155572

85. Wu, J., Hu, R., Li, D., (...), Zang, Y. (2024) A GNN-based fraud detector with dual resistance to graph disassortativity and imbalance. *Information Sciences*. https://doi.org/10.1016/j.ins.2024.120580

86. Chakraborty, N.R., Chowdhury, M.J.M., Rajaeian, M., (...), Biswas, K. (2026) A Comprehensive Evaluation of GNN Models for Fraud Detection in Blockchain Transactions. *2026 IEEE International Conference on Blockchain and Cryptocurrency, ICBC 2026*. https://doi.org/10.1109/ICBC67748.2026.11575502

87. Liu, Z., Dou, Y., Yu, P.S., (...), Peng, H. (2020) Alleviating the Inconsistency Problem of Applying Graph Neural Network to Fraud Detection. *SIGIR 2020 - Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*. https://doi.org/10.1145/3397271.3401253